

We thank all reviewers for their constructive comments! Please find our responses below.

More Mediator results Lora merging (by reviewer Reviewer #aCDT) and AdaMerging and WeMOE (by reviewer #Q1B5).

Table 1. Comparing performance of different model merging methods on Qwen 1.5 4B model. We added LoRA-Merging for Reviewer aCDT, AdaMerging and WeMOE for Reviewer Q1B5 to provide comprehensive comparisons.

Algorithm/Tasks	GSM8K	TriviaQA	Winogrande	HumanEval	MMLU	All Tasks Average
base	47.16	44.54	56.75	41.46	54.45	48.87
Math	51.00	46.95	54.62	26.83	53.54	46.79
Code	43.29	46.39	54.14	43.29	54.82	48.39
QA	45.56	48.02	57.93	39.02	52.32	48.57
all-sft	48.52	47.73	55.88	39.14	53.93	49.04
TIES	47.76	46.59	54.14	44.51	54.58	49.5
AdaMerging	47.46	47.90	54.38	44.51	54.62	49.59
PCB-merging	47.83	47.60	56.75	43.90	54.58	49.93
Twin-merging	47.99	44.63	57.54	40.85	52.98	48.80
LoRA-Merging	48.02	44.69	57.61	40.85	53.01	48.87
BTX	48.44	46.94	57.77	42.68	53.88	49.94
WeMOE	47.83	47.84	53.99	45.73	54.70	49.83
Mediator	50.94	48.20	57.85	45.12	54.87	51.40

- **Lora-Merging (by reviewer Reviewer #aCDT).** Table 1 show that while LoRA merging slightly outperforms Twin-merging, both matrix decomposition approaches introduce additional approximation errors. In contrast, **our merging method achieves better performance by avoiding these decomposition errors and preserving the original parameter values more faithfully.
- **AdaMerging and WeMOE (by reviewer #Q1B5).** We also add experiments to compare with WEMoE and Adamerging. The following table shows that Mediator achieved the best performance through global task-level routing and enhanced classification precision that ensures accurate expert selection for non-OOD tasks; Similar performance gaps are also observed when comparing with AdaMerging and PCB-Merging.

Mediator routing results (by reviewer Reviewer #aCDT)

The table shows the total number of layers requiring routing for each model architecture:

- Qwen-4B requires routing in 6 layers total:
 - 5 layers in lower positions (layers 0-4)
 - 1 layer in higher position (layer 39)
- Qwen-7B requires routing in 7 layers total:
 - 6 layers in lower positions (layers 0-5 and 9)
 - 1 layer in higher position (layer 27)
- LLaMA-3.2-3B requires routing in 5 layers total:
 - 4 layers in lower positions (layers 0,1,2,5)
 - 1 layer in higher position (layer 31)
- LLaMA-3.1-8B requires routing in 7 layers total:
 - All 7 routing layers are in lower positions (layers 0-6)
 - No routing needed in higher layers

This analysis shows that across different model architectures, the number of layers requiring routing remains relatively consistent (5-7 layers), with most routing concentrated in the lower layers.

Mediator routing results (by Reviewer #aCDT)

Answer Our analysis reveals that Mediator employs an efficient layer-wise routing strategy that enables significant memory savings while preserving model performance. The key findings are:

1) Adaptive Layer Selection: Mediator automatically determines which layers require routing based on parameter conflict ratios, applying routing only when conflicts exceed one standard deviation from the mean. This data-driven approach eliminates the need for manual hyperparameter tuning.

2) Minimal Routing Requirements: Across architectures, only a small fraction of layers need routing: - Qwen-4B: 15% (6/40 layers) - Qwen-7B: 25% (7/28 layers) - LLaMA-3.2-3B: 15.6% (5/32 layers) - LLaMA-3.1-8B: 25% (7/28 layers)

3) Strategic Layer Distribution: As shown in Table 2, routing is concentrated in the lower layers where task-specific adaptations are most critical, with selective routing in higher layers based on conflict analysis.

This optimized approach enables Mediator to achieve over 4x memory compression compared to storing separate experts, while maintaining model quality through targeted parameter sharing and routing. The layer-wise strategy emerges naturally from analyzing parameter conflicts rather than requiring manual tuning.

Table 2. Layer-wise routing strategy for different model architectures

Models	Lower routing layers	Higher routing layers
Qwen-4B	0-4	39
Qwen-7B	0-5, 9	27
llama-3.2-3B	0,1,2,5	31
llama-3.1-8B	0-6	NA

Question: Better illustration of Fig. 3 in the paper (by Reviewer #sV1X).

Question: More insight in COT experiments (by Reviewer #sV1X).

Following the suggestions of reviewer #sV1X, we redraw the Fig 3 as shown in Fig. 1 as follows.

Table 3 and Table 4 show the detailed results of our experiments on these challenging benchmarks and also compares the results between SFT models on COT and non-cot datasets.

Our ablation studies demonstrate the domain-specific benefits of Chain-of-Thought (CoT) prompting:

- In Table 4, removing Math CoT mainly affects math tasks (GSM8K: -2.56%, GSM-Plus-mini: -2.26%) .
- Similarly, removing Code CoT primarily impacts coding tasks (HumanEval: -0.61%, MBPP: -1.00%) .

On the challenging benchmarks, Mediator maintains strong performance:

- For GSM-Plus-mini, Mediator (36.25%) outperforms the base model by 3.42% and leads other merging methods.
- On MBPP, Mediator achieves 34.20%, showing a 1.00% improvement over the base model.

These results validate our COT finetuning methods effectiveness on complex reasoning tasks.

Question: GPT-4 results of tasks (by reviewer #Q1B5). We presents the results of GPT-4 across multiple tasks, as shown in Table . The results demonstrate strong performance across diverse tasks, with GPT-4 achieving over 87% accuracy on all benchmarks. Particularly notable is the 94.8% accuracy on GSM8K, showing exceptional mathematical reasoning capabilities. The consistent high performance across different task types - from mathematical reasoning (GSM8K) to code generation (HumanEval) and general knowledge (MMLU, TriviaQA).

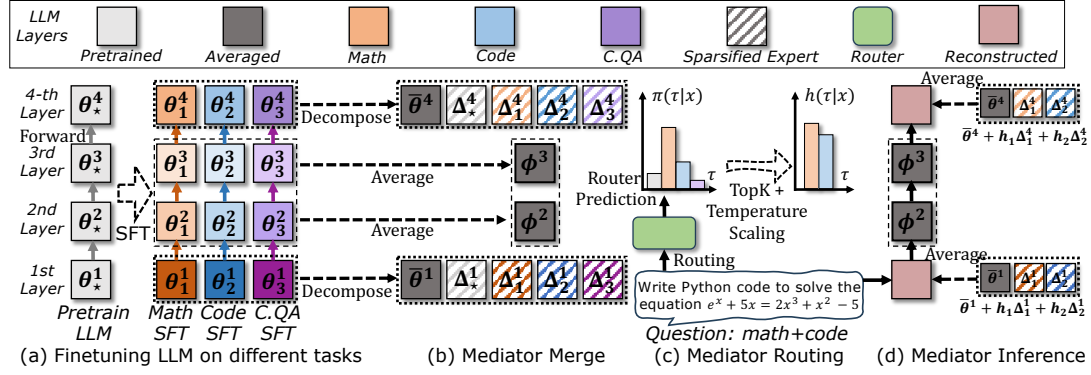


Figure 1. The framework of Mediator. For three finetuned LLMs on different tasks, Mediator decomposes layers with large conflicts into expert task arithmetics and sparsifies them. And Mediator averages layers with less conflicts. During inference, Mediator inputs the input question to the router, which output the logits of different experts. Then, top-K experts are chosen and averaged with the base model with temperature scaling weights. Then, the input question is sent to the final merged model.

Table 3. Results on additional challenging reasoning benchmarks

Algorithm/Tasks	GSM-Plus-mini	MBPP
Base	32.83	33.20
Math Expert	36.71	32.8
Code Expert	34.12	34.2
TIES	34.12	34.00
Twin-merging	34.12	34.00
PCB-merging	34.12	34.20
BTX	35.79	34.00
Mediator	36.25	34.20

Table 4. Ablation study w/o CoT (format: w/o CoT/with CoT, ↑ indicates improvement)

Algorithm/Tasks	GSM8K	HumanEval	GSM-Plus-mini	MBPP
w/o Math CoT	48.44/51.00 (↑2.56)	27.44/26.83	34.45/36.71 (↑2.26)	33.20/32.8
w/o Code CoT	46.95/43.29	42.68/43.29 (↑0.61)	33.29/34.12 (↑0.83)	33.20/34.20 (↑1.00)

Task	Performance
GSM8K	94.8
TriviaQA	87.0
WinoGrande	87.5
HumanEval	90.2
MMLU	87.0

Table 5. GPT-4 Performance Across Different Tasks